

LoadGen

Goals for the creation of this utility was to generate loads against REST endpoints and then to be able to describe a platform profile. There generally seems to be two types of load testing.

1. Vendor testing:
 - vehicle's top speed
 - top ops per second under specific conditions (i.e. 100 million entries with xxxxx authentications per second across two OpenAM instance with 4 cores and 16GB memory each)
2. Customer testing/platform profile
 - driving this vehicle with this load over this type of terrain under these conditions **should** yield a fuel consumption rate of X.
 - A 4 core system with 16GB of memory with 100 million entries and a 10% concurrency rate performing these types of task(s), for example:
 1. Authenticate to OpenAM
 2. sleep 10 seconds
 3. Validate session
 4. sleep 10 seconds
 5. Update user data
 6. Logout**should** provide a transaction rate of X with a 200ms or less latency.

Why I put the effort into this? Scripted cURL (UNIX shell) is resource intensive on the load generator. Used REST so eventually loads can be applied to competitive offerings using the same utility and makes testing ForgeRock services much easier and flexible. Looked at Jmeter and gatling.io. Very good utilities but the learning curve to achieve the desired results seemed high for both. Plus needed a utility that helped craft a profile of the platform not just stress the platform.

LoadGen uses REST to interact with the services (OpenAM, OpenDJ, OpenIDM) and uses json for configuration and testing parameters.

Job example (\$LOADGEN_HOME/config):

```
{
  "name": "my job 0",
  "service-environment": "one 4xcore 7GB Ubuntu 14.04 VM in Azure Apache Tomcat 7.59, JDK8, OpenAM 12.0.0 with 200,000 users",
  "service-location-port": "http://openam.example.com:8080",
  "ADMINUSERID": "adminuser",
  "ADMINUSERPASSWORD": "password",
  "load-generator": "one 2xcore Ubuntu 14.04 VM in Azure on same vnet at OpenAM server (example.com)",
  "threads" :20,
  "threadgroupsize" : 2,
  "threadinterval" : 10000,
  "upperbounds" : 10000,
  "workloads": [
    {
      "read-user-object" : 1000
    },
    {
      "query-user-objects" : 100
    },
    {
      "authenticate-read-logout-user-object" : 0
    },
    {
      "create-and-update-user-object": 20
    }
  ]
}
```

The value after each workload such as “read-user-object” is the number of iterations that the workload should be performed. The **threads** value is the number of Java threads that will be created and executed. Each thread will run the same workloads in the same order. Workloads are intermixed based on the quantity of the various workloads. If there are 10 workloads of type A and 2 workloads of type B then Workload A will be executed 5 times, followed by 1 Workload B, then 5 more Workload A and 1 Workload B.

The **threadgroupsize** and **threadinterval** help regulate the initial spawning of threads. Helps with creating a ramped up load and sustainable load. Set to zero to turn off and have all threads start at the same time.

Some attributes/values can be defined at all three levels - jobs, workloads, tasks – such as: service-location, service-port, adminuserid, adminuserpassword, randomfile, etc. Precedence: tasks override workloads; workloads override jobs.

workload example (\$LOADGEN_HOME/workloads):

```
{
  "name": "authenticate-read-logout-user-object",
  "persist" : [tokenId, ...]
  "tasks": [
    {
      "authenticate-openam-user-object": 400
    },
    {
      "sleep" : 10
    },
    {
      "read-openam-user-object" : 400
    },
    {
      "logout-openam-user-object" : 100
    }
  ]
}
```

The value after each task is the millisecond threshold that the task must complete in (except for “sleep”). If the task is not completed in under the specified ms then it is counted as “exceeded”. If the task is “sleep” than the value is added to the millisecond value specified in the sleep's “url-payload”. Using a ramped up load generation (using threadgroupsize and threadinterval) can skew certain aspects of the results. For example total successful operations should be considered invalid when a ramped up load is utilized.

task examples; three are listed below. Essentially the data for a operational cURL command. (\$LOADGEN_HOME/tasks):

```
{
  "name": "authenticate-openam-user-object",
  "service-location-port" : "http://openam0.example.com:8080",
  "request": "POST",
  "header": {
    "X-OpenAM-Username": "user.$RANDOMVALUE",
    "X-OpenAM-Password": "password",
    "Content-Type": "application/json"
  },
  "url-endpoint": "/openam/json/",
  "url-payload": "authenticate",
  "data-payload": "{}"
}
-----
{
  "name": "create-opensj-user-object",
  "request": "PUT",
  "header": {
    "X-OpenIDM-Username": "$ADMINUSERID",
    "X-OpenIDM-Password": "$ADMINUSERPASSWORD",
    "Content-Type": "application/json",
    "If-None-Match": "*"
  },
  "url-endpoint" : "/users/",
  "url-payload": "user.new.$RANDOMVALUE",
  "data-payload": {
    "_id": "user.$RANDOMVALUE",
    "contactInformation": {
      "telephoneNumber": "+1 408 555 1212",
      "emailAddress": "user.new.$RANDOMVALUE@example.com"
    },
    "name": {
      "familyName": "New.$RANDOMVALUE",
      "givenName": "User"
    },
    "displayName": "User New $RANDOMVALUE"
  }
}
```

```
}  
}  
-----  
{  
  "name": "sleep",  
  "service-location-port" : "http://localhost:8080",  
  "header": {  
    "Content-Type": "application/json"  
  },  
  "url-endpoint": "$SLEEP",  
  "url-payload": 100  
}
```

Output from a job with no ramp up:

Job Started with 4 threads on Fri Jun 26 11:58:36 PDT 2015...and completed Fri Jun 26 11:59:43 PDT 2015

| Operation | TxTotal | AccmTime | Thrshold | TxPass | PassTime | TxExced | ExcdTime | TxFail | Failtime | CbdPsOps | ThrdOps | Avrms/op | Success | Exceed | Fail |
|-------------------------------|---------|----------|----------|--------|----------|---------|----------|--------|----------|-----------|-----------|-------------|---------|--------|-------|
| read-opendj-user-object | 40000 | 211.088s | 400ms | 39991 | 211.023s | 0 | 0.000s | 9 | 0.065s | 758.041/s | 189.510/s | 5.277ms/op | 99.98% | 0.00% | 0.02% |
| query-opendj-user-objects | 4000 | 34.814s | 400ms | 3999 | 34.809s | 0 | 0.000s | 1 | 0.005s | 459.536/s | 114.884/s | 8.704ms/op | 99.97% | 0.00% | 0.03% |
| create-opendj-user-object | 40 | 1.803s | 400ms | 40 | 1.803s | 0 | 0.000s | 0 | 0.000s | 88.741/s | 22.185/s | 45.075ms/op | 100.00% | 0.00% | 0.00% |
| update-opendj-new-user-object | 40 | 1.032s | 200ms | 38 | 0.596s | 2 | 0.436s | 0 | 0.000s | 255.034/s | 63.758/s | 15.684ms/op | 95.00% | 5.00% | 0.00% |
| read-opendj-new-user-object | 80 | 1.277s | 200ms | 76 | 1.277s | 0 | 0.000s | 4 | 0.000s | 238.058/s | 59.514/s | 16.803ms/op | 95.00% | 0.00% | 5.00% |

Job lapsed time = 66399ms

Passed Accum time = 249508ms

Total Threads = 4

Thread group size = 0

Launch interval = 1000ms

Transaction goal = 44160

Transction actual = 44144

Successful ops/s = 707.697

Ops/s during job = 664.829

Success rate = 99.96%

Total sleep sec = 12.179

Based on following workloads:

Workload: read-user-object with 40000 loops; Percent load = 90.33%

Tasks used:

read-opendj-user-object

Workload: query-user-objects with 4000 loops; Percent load = 9.03%

Tasks used:

query-opendj-user-objects

Workload: create-and-update-user-object with 40 loops; Percent load = 0.63%

Tasks used:

create-opendj-user-object

sleep

update-opendj-new-user-object

sleep

read-opendj-new-user-object

sleep

read-opendj-new-user-object

Target environment(s): mac laptop running two OpenDJ 2.6.2 instance in vbox, replicated, behind haproxy

Load generator: mac laptop running JDK 8

Some preliminary findings:
Increased thread count quickly reduces the transaction rate per thread. While two threads may scale up against a two core system two threads do NOT see double the performance over one thread. In fact OpenAM, with a simple authentication/validation workload did not scale to 4 cores as well as hoped.

For example a 4 core OpenAM instance with 16 threads:

=====

Job Started with 16 threads on Thu Jun 25 15:43:06 UTC 2015...and completed Thu Jun 25 15:55:17 UTC 2015

| Operation | TxTotal | AccmTime | Thrshold | TxPass | PassTime | TxExced | ExcdTime | TxFail | Failtime | CbdPsOps | ThrdOps | Avrms/op | Success | Exceed | Fail |
|---------------------------------|---------|------------|----------|--------|------------|---------|----------|--------|----------|------------|-----------|-------------|---------|--------|-------|
| authenticate-openam-user-object | 160000 | 11130.598s | 200ms | 159679 | 11065.405s | 321 | 65.193s | 0 | 0.000s | 230.888/s | 14.430/s | 69.298ms/op | 99.80% | 0.20% | 0.00% |
| logout-openam-user-object | 160000 | 545.667s | 100ms | 159639 | 541.514s | 40 | 4.153s | 321 | 0.000s | 4716.820/s | 294.801/s | 3.392ms/op | 99.77% | 0.03% | 0.20% |

Job lapsed time = 731276ms

Passed Accum time = 11606919ms

Total Threads = 16

Transaction goal = 320000

Transction actual = 319318

Successful ops/s = 440.176

Ops/s during job = 436.659

Success rate = 99.79%

Total sleep sec = 0.000

Based on following workloads:

Workload: authenticate-read-logout-user-object with 160000 loops; Percent load = 100.00%

Tasks used:

authenticate-openam-user-object

logout-openam-user-object

Target environment(s): one 4xcore 7GB Ubuntu 14.04 VM in Azure Apache Tomcat 7.59, JDK8, OpenAM 12.0.0 with 200,000 users

Load generator: one 2xcore Ubuntu 14.04 VM in Azure on same vnet at OpenAM server (example.com)

=====

And with a 8 core OpenAM instance with 16 threads

=====

Job Started with 16 threads on Thu Jun 25 19:54:52 UTC 2015...and completed Thu Jun 25 20:02:00 UTC 2015

| Operation | TxTotal | AccmTime | Thrshold | TxPass | PassTime | TxExced | ExcdTime | TxFail | Failtime | CbdPsOps | ThrdOps | Avrms/op | Success | Exceed | Fail |
|---------------------------------|---------|-----------|----------|--------|-----------|---------|----------|--------|----------|------------|-----------|-------------|---------|--------|-------|
| authenticate-openam-user-object | 160000 | 6389.917s | 200ms | 159632 | 6313.170s | 368 | 76.747s | 0 | 0.000s | 404.569/s | 25.286/s | 39.548ms/op | 99.77% | 0.23% | 0.00% |
| logout-openam-user-object | 160000 | 443.919s | 100ms | 159632 | 443.919s | 0 | 0.000s | 368 | 0.000s | 5753.554/s | 359.597/s | 2.781ms/op | 99.77% | 0.00% | 0.23% |

Job lapsed time = 428219ms

Passed Accum time = 6757089ms

Total Threads = 16

Transaction goal = 320000

Transction actual = 319264

Successful ops/s = 755.980

Ops/s during job = 745.562

Success rate = 99.77%

Total sleep sec = 0.000

Based on following workloads:

Workload: authenticate-read-logout-user-object with 160000 loops; Percent load = 100.00%

Tasks used:

authenticate-openam-user-object

logout-openam-user-object

Target environment(s): one 8xcore 14GB Ubuntu 14.04 VM in Azure Apache Tomcat 7.59, JDK8, OpenAM 12.0.0 with 200,000 users

Load generator: one 2xcore Ubuntu 14.04 VM in Azure on same vnet at OpenAM server (example.com)

=====

It will be very important to understand the load breakout. Does the potential workload include many threads (1,000 – 100,000) that are somewhat concurrent against a service endpoint or a concentrated set of threads (100 or less)? For example mobile devices hitting an endpoint directly would constitute a very high thread count which may be best served by many small (2 core) OpenAM servers. Large quantities of threads seem to cause as much stress as the type of workload they represent. It seems to be just as important to understand the workload as how the workload is delivered.

When a system started to stress the number of operations exceeding threshold started to rise quickly and “Successful ops/s” started to drop. Thresholds can be adjusted higher but understanding how quickly a task can be responded to is very important. For example OpenDJ reads can occur in under 100ms very easily (from the load generators standpoint). Adjusting the threshold to something higher may result in a higher transaction rate but in fact if OpenDJ reads start taking longer than 50ms (from the load generators standpoint) the service is in fact starting to show stress. OpenDJ updates on the other hand can exceed 400ms regularly but the service is not stressed. Best way to establish a baseline threshold (one where the service is not stressing and consistently delivering expected performance) is to run one thread against a service with at least two cores, plenty of memory and is properly tuned,

for an extended amount of time (one hour). This “unstressed” service should be handle to respond consistently. The “Avrms/ops” data reflects the average time, in milliseconds, the load generator “waited” for the service to complete the transaction.

Generally, when a customer asks for sizing information, they want to put in place a system that will handle the average load without undue stress. This way if a spike occurs there is headroom for the service to handle the spike without impacting performance (the end user's experience for example).

With simple tests is was next to impossible to drive the service to 100% busy from a CPU standpoint (seemed true for both OpenDj and OpenAM with the very basic testing thus far).

A multi-core load generator was able to push more load against the service than multiple load generating instances (separate Vms). For example a 4 core load generator can kick out more transactions than two 2 core load generators.

Additional capability that is needed:

1. Website/REST endpoint to provide sleep as well as OpenAM agent
2. Random reads and injection from files for user ids and data
3. Attributes definable at jobs, workloads, and tasks levels with said precedence.
4. Website to build and run tests. Test as a service or trial or download war from backstage (need to authenticate)
5. Fix PATCH for OpenDJ and OpenIDM
6. Workload needs to specify if value needs to be persisted across tasks by specifying the attribute name returned by the a task operation i.e. tokenId
7. Define wildcards:

```
WILDCARDS : [  
  {  
    name : $RANDOMVALUE,  
    type : generated  
  },  
  {  
    name : $TOKENID,  
    type : retained  
  },  
  {  
    name : $FILE-A,  
    type : file,  
    source : path to file  
  },  
  {  
    name : $OBJECT-O,  
    type : rest-endpoint,  
    source : url to object  
  }  
]  
}
```

Sample output as of October 19th 2016:
rmfaller-forgerock-mac:LoadGen rmfaller\$ java -jar ./dist/LoadGen.jar --config ./jobs/jobs-openam.json
Job Started with 4 threads on Wed Oct 19 15:37:24 PDT 2016...and completed Wed Oct 19 15:37:25 PDT 2016

| Operation | TxTotal | AccmTime | Thrshold | TxPass | PassTime | TxExced | ExcdTime | TxFail | Failtime | CbdPsOps | ThrdOps | Avrms/op | Success | Exceed | Fail |
|--------------------------------|---------|----------|----------|--------|----------|---------|----------|--------|----------|-----------|----------|-------------|---------|--------|-------|
| authenticate-openam-admin | 40 | 0.816s | 400ms | 40 | 0.816s | 0 | 0.000s | 0 | 0.000s | 196.078/s | 49.020/s | 20.400ms/op | 100.00% | 0.00% | 0.00% |
| create-openam-user-object-post | 40 | 1.015s | 800ms | 40 | 1.015s | 0 | 0.000s | 0 | 0.000s | 157.635/s | 39.409/s | 25.375ms/op | 100.00% | 0.00% | 0.00% |
| logout-openam-admin | 40 | 0.440s | 100ms | 40 | 0.440s | 0 | 0.000s | 0 | 0.000s | 363.636/s | 90.909/s | 11.000ms/op | 100.00% | 0.00% | 0.00% |

Job lapsed time = 617ms
Passed Accum time = 2271ms
Total Threads = 4
Thread group size = 0
Launch interval = 0ms
Transaction goal = 120
Transaction actual = 120
Successful ops/s = 211.361
Ops/s during job = 194.489
Success rate = 100.00%
Total sleep sec = 0.000

Based on following workloads:
 Workload: adminuser-authenticate-createuser-logout with 40 loops; Percent load = 100.00%
 Tasks used:
 authenticate-openam-admin
 create-openam-user-object-post
 logout-openam-admin

Target environment(s): mac laptop running vbox hosting ubuntu (2 cores, 8GB) running an OpenAM 13.5 instance loaded w/2000 objects
Load generator: mac laptop running JDK 8

| | |
|-----------|---|
| Operation | from the ./tasks JSON formatted REST call |
| TxTotal | number of actual tasks executed (threads * workload count {as specified in the ./jobs description file}) |
| AccmTime | accumulated time in seconds to execute all Operations (successful and unsuccessful) |
| Thrshold | maximum time in milliseconds allowed for the Operation to complete successfully and specified in the ./workloads description file |
| TxPass | number of Operations that passed successfully (did not timeout or incur any other type of error) |
| PassTime | accumulated time in seconds to execute all successful Operations |
| TxExced | number of Operations that exceeded the specified Threshold |
| ExcdTime | accumulated time in seconds spent on Operations that eventually exceeded the Threshold |
| TxFail | number of Operations that failed for issues outside of exceeding the specified Threshold |
| Failtime | accumulated time in seconds spent on Operations that failed for issues outside of exceeding the specified Threshold |
| CbPsOps | calculated number of successful Operations that could be executed by all threads in a second |
| ThrdOps | calculated number of successful Operations per thread that could be executed in a second |
| Avrms/op | average number of milliseconds needed to execute one Operation |
| Success | percentage of successful Operations (this value degrades further if following Operations are dependent on the success of previous Operations) |
| Exceed | percentage of Operations that exceeded the threshold |
| Fail | percnetage of Operations that failed for reasons other than exceeding the threshold |

| | |
|-------------------|--|
| Job lapsed time | clock time in milliseconds it took for the job to complete |
| Passed Accum time | total accumulated time of all successful Operations |
| Total Threads | as specified in the ./jobs file |
| Thread group size | used to spread load over time; of the total number of Threads how many threads to start together |

| | |
|---------------------------|--|
| Job lapsed time | clock time in milliseconds it took for the job to complete |
| Launch interval | time in milliseconds to wait before executing a thread group |
| Transaction goal | total number of transactions that will hopefully complete without exceeding the thresholds or an error |
| Transaction actual | reality |
| Successful ops/s | For the entire clock time of the job the number of Operations per second that completed successfully |
| Ops/s during job | calculated average Operations per second based on completion time of each Operation |
| Success rate | overall total success rate of all Operations for the job |
| Total sleep sec | if specified the total time in seconds that threads spent sleeping |